PROJECT TITLE[1]

Author name 1
Affiliation

Author name 2
Affiliation

Author name 3
Affiliation

Author name 4
Affiliation

Author name 5
Affiliation

Author name 6
Affiliation

With
Apart Research

Abstract

Summarize your project in 150–250 words. A strong abstract lets a reviewer understand what you did and why it matters without reading anything else. Make sure to cover: the problem, your approach, key results, and the main takeaway. Polish it last: the abstract should reflect your final results, not your initial plan.

How to use this template: Replace the *italicized guidance text* under each section with your content. The section structure is strong guidance but not rigid. If your project requires a different organization, feel free to adapt. Delete all guidance text including this info box before submitting. Make sure the project title and author information above are up to date.

How your submission will be evaluated: Your project will be judged on the quality of this written report. To understand what we look for, see our Evaluation Rubric.

Recommended length: 4 pages excluding references and appendix.
Rough guide: Intro & Related Work 1p, Methods and Results: 2.5p, Discussion 0.5p.

[1] Research conducted at the The Global South AI Safety Hackathon, June 2026

1. Introduction

What problem are you addressing and why does it matter?
Connect it to your work: we want to know why your work is practically valuable.

Provide enough background for readers to understand your work.
If relevant, briefly describe the threat model or failure mode you're addressing; reference prior work that motivates why it is worth addressing, or explain it yourself.

Aspire to clearly list your most important contributions that go beyond what exists today.
Our main contributions are:

1. *[First contribution — what new thing did you create, discover, or demonstrate?]*
2. *[Second contribution]*
3. *[Third contribution, if applicable]*

2. Related Work

What prior work is most similar, and how does your work differ?
Cite the most relevant papers, tools, or projects. Explain what gap your work addresses.

Some questions which may help:
- *(a) When and why would someone use your method over the existing state-of-the-art?*
- *(b) What information/insight does your method provide which we did not have before?*

3. Methods

Describe your approach clearly enough that someone could replicate it. Include key design choices and justify them where relevant (hint: the more you can back up your design choices by referencing prior work, the better).

What models/datasets/tools did you use and why? What were key parameters or design decisions? What did you try that didn't work? Could someone reproduce your work from this description?

4. Results

Present your main findings with appropriate evidence. Use figures and tables where appropriate (we strongly encourage at least one figure, see tips below). Distinguish between observations and interpretations.

Argue why your claims are robust. E.g., if your approach "performs better" than alternatives:

(a) *Do you have enough data? Is the difference statistically significant?*
(b) *Is it robust? Or do small changes to your setup cause substantial changes to your results?*

Tips for figures and tables:

- *Number all figures (Figure 1, Figure 2...) and tables (Table 1, Table 2...)*
- *Include descriptive captions that can be understood without the main text*
- *Place figures/tables near where they're first referenced*
- ***Ensure text in figures is legible!***

5. Discussion and Limitations

Discuss the broader implications for AI safety.
What do your results mean? What trends do you notice and what might they indicate?

Limitations

What are the limitations of your work? What threat models or failure modes did you not address? Be honest about constraints — methodological limitations, scope limitations, or aspects you couldn't fully address in the hackathon timeframe. Explicitly note the assumptions you made, whether implicitly or explicitly, and how the interpretation of your results would change if a given assumption did not hold.

Future Work

What are the natural next steps? How could this work be extended?

6. Conclusion

Briefly summarize your main findings and their implications (1–2 paragraphs).

Code and Data

Include links if applicable. If your project doesn't involve code (e.g., policy analysis) or if there are info-hazard considerations, note that here.

- ***Code repository****: [Link to GitHub/GitLab if applicable]*
- ***Data/Datasets****: [Link if applicable]*
- ***Other artifacts*** *(optional): [Demo link, video walkthrough, Hugging Face Space, etc.]*

Author Contributions (optional)

[e.g., "A.B. led the project and designed experiments. C.D. implemented the code. All authors contributed to writing and reviewed the final manuscript."]

References

Use a consistent citation format. Include: Author(s), Year, Title, Venue/Publisher, and URL or DOI where available.

1. *[Reference 1]*
2. *[Reference 2]*
3. ...

Appendix *(optional)*

Supplementary material such as additional figures, detailed methodology, prompts used, extended results, etc.

LLM Usage Statement

If you used LLM assistance in developing your project or writing this report, briefly note how. Ensure all claims and results have been verified.

NOTE: We strongly encourage that the final version of the submission is primarily written by your team.

[e.g., "We used Claude to brainstorm approaches and help draft sections. All results and claims were independently verified."]